

## NPFL114, Lecture 10

# Sequence Prediction, Reinforcement Learning



Milan Straka

## Structured Prediction

# Structured Prediction



Consider generating a sequence of  $y_1, \dots, y_N \in Y^N$  given input  $\mathbf{x}_1, \dots, \mathbf{x}_N$ .

# Structured Prediction



Consider generating a sequence of  $y_1, \dots, y_N \in Y^N$  given input  $\mathbf{x}_1, \dots, \mathbf{x}_N$ . Predicting each sequence element models the distribution  $P(y_i | \mathbf{X})$ .

# Structured Prediction




Consider generating a sequence of  $y_1, \dots, y_N \in Y^N$  given input  $\mathbf{x}_1, \dots, \mathbf{x}_N$ . Predicting each sequence element models the distribution  $P(y_i | \mathbf{X})$ . There are two problems with the above approach:

1. There may be dependencies among the  $y_i$  themselves.

# Structured Prediction

Consider generating a sequence of  $y_1, \dots, y_N \in Y^N$  given input  $\mathbf{x}_1, \dots, \mathbf{x}_N$ . Predicting each sequence element models the distribution  $P(y_i | \mathbf{X})$ . There are two problems with the above approach:

1. There may be dependencies among the  $y_i$  themselves.
2. Even if we consider dependencies, if we compute softmax for each  $y_i$  individually, we will hit the *label bias problem*.



# Seq2seq Beam Search Decoding



So far we described only greedy decoding in a sequence to sequence decoder.

However, such decoding might not find the most probable output sequence.

We might consider a *beam search*, where we iteratively compute some fixed number (a *beam size*) of best output sequences.

# Conditional Random Fields



Let  $G = (V, E)$  be a graph such that  $Y$  is indexed by vertices of  $G$ . Then  $(\mathbf{X}, \mathbf{y})$  is a conditional markov field, if the random variables  $\mathbf{y}$  conditioned on  $\mathbf{X}$  obey the Markov property with respect to the graph, i.e.,

$$P(y_i | \mathbf{X}, y_j, i \neq j) = P(y_i | \mathbf{X}, y_j, (i, j) \in E).$$



# Conditional Random Fields



Let  $G = (V, E)$  be a graph such that  $Y$  is indexed by vertices of  $G$ . Then  $(\mathbf{X}, \mathbf{y})$  is a conditional markov field, if the random variables  $\mathbf{y}$  conditioned on  $\mathbf{X}$  obey the Markov property with respect to the graph, i.e.,

$$P(y_i | \mathbf{X}, y_j, i \neq j) = P(y_i | \mathbf{X}, y_j, (i, j) \in E).$$

Usually we assume that dependencies of  $\mathbf{y}$ , conditioned on  $\mathbf{X}$ , form a chain.

# CRF Output Layer



For a sequence of  $(\mathbf{x}_1, \dots, \mathbf{x}_N)$  and  $(y_1, \dots, y_N)$ , we define a score as

$$s(\mathbf{X}, \mathbf{y}; \boldsymbol{\theta}, \mathbf{A}) = \sum_{i=1}^N (\mathbf{A}_{y_{i-1}, y_i} + f_{\boldsymbol{\theta}}(y_i | \mathbf{X})).$$

# CRF Output Layer

For a sequence of  $(\mathbf{x}_1, \dots, \mathbf{x}_N)$  and  $(y_1, \dots, y_N)$ , we define a score as

$$s(\mathbf{X}, \mathbf{y}; \boldsymbol{\theta}, \mathbf{A}) = \sum_{i=1}^N (\mathbf{A}_{y_{i-1}, y_i} + f_{\boldsymbol{\theta}}(y_i | \mathbf{X})).$$

We then define

$$p(\mathbf{y} | \mathbf{X}) = \text{softmax}_{\mathbf{z} \in Y^N} (s(\mathbf{X}, \mathbf{z}))_{\mathbf{z}},$$

so that

$$\log p(\mathbf{y} | \mathbf{X}) = s(\mathbf{X}, \mathbf{y}) - \text{logadd}_{\mathbf{z} \in Y^N} (s(\mathbf{X}, \mathbf{z})).$$

# CRF Output Layer



We can compute  $p(\mathbf{y}|\mathbf{X})$  efficiently using dynamic programming. If we denote  $\alpha_t(k)$  as probability of all sentences with  $t$  elements with the last  $y$  being  $k$ .

We can then show that

$$\begin{aligned}\alpha_t(k) &= \text{logadd}_z(s(\mathbf{X}, \mathbf{z})) \\ &= f_{\theta}(y_t = k | \mathbf{X}_{1:t}) + \text{logadd}_i(\alpha_{t-1}(i) + \mathbf{A}_{i,k}).\end{aligned}$$

For efficient implementation, we use the fact that

$$\ln(a + b) = \ln a + \ln(1 + e^{\ln b - \ln a}).$$


# Connectionist Temporal Classification



Let us again consider generating a sequence of  $y_1, \dots, y_M$  given input  $\mathbf{x}_1, \dots, \mathbf{x}_N$ , but this time  $M \leq N$  and there is no explicit alignment of  $\mathbf{x}$  and  $y$  in the gold data.

# Connectionist Temporal Classification

Let us again consider generating a sequence of  $y_1, \dots, y_M$  given input  $\mathbf{x}_1, \dots, \mathbf{x}_N$ , but this time  $M \leq N$  and there is no explicit alignment of  $\mathbf{x}$  and  $y$  in the gold data.



# Connectionist Temporal Classification



We enlarge the set of output labels by a – (*blank*) and perform a classification for every input element to produce an *extended labeling*. We then post-process it by the following rules:

1. We remove neighboring symbols.
2. We remove the –.

# Connectionist Temporal Classification



We enlarge the set of output labels by a – (*blank*) and perform a classification for every input element to produce an *extended labeling*. We then post-process it by the following rules:

1. We remove neighboring symbols.
2. We remove the –.

Because the explicit alignment of inputs and labels is not known, we consider *all possible* alignments.



# Connectionist Temporal Classification



Let us denote the probability of label  $l$  at time  $t$  as  $p_l^t$ .

We now consider a modified label sequence by inserting a – to the beginning, to the end and between every pair of labels. Using this modified labeling we define

$$\alpha_t(s) \stackrel{\text{def}}{=} \sum_{\text{labeling } \pi \text{ with labels of } \pi_{1:t} = \mathbf{y}_{1:s}} \prod_{t'=1}^t p_{\pi_{t'}}^{t'}.$$


# Connectionist Temporal Classification

Let us denote the probability of label  $l$  at time  $t$  as  $p_l^t$ .

We now consider a modified label sequence by inserting a – to the beginning, to the end and between every pair of labels. Using this modified labeling we define


$$\alpha_t(s) \stackrel{\text{def}}{=} \sum_{\text{labeling } \pi \text{ with labels of } \pi_{1:t} = \mathbf{y}_{1:s}} \prod_{t'=1}^t p_{\pi_{t'}}^{t'}.$$

Similarly to the CRF, we can use dynamic programming to compute  $\alpha$  in polynomial time.



# CTC Decoding

Unlike CRF, we cannot easily perform the decoding in an optimal way. We therefore either use greedy decoding, or a beam search.




$$\begin{aligned} p(l=\text{blank}) &= p(- -) \\ &= 0.7 * 0.6 \\ &= 0.42 \end{aligned}$$


$$\begin{aligned} p(l=A) &= p(AA) + p(A-) + p(-A) \\ &= 0.3 * 0.4 + 0.3 * 0.6 + 0.7 * 0.4 \\ &= 0.58 \end{aligned}$$

## Reinforcement Learning

# Reinforcement Learning




# Reinforcement Learning



A Markov decision process is a quadruple  $(\mathcal{S}, \mathcal{A}, P, \gamma)$ , where:

- $\mathcal{S}$  is a set of states,
- $\mathcal{A}$  is a set of actions,
- $P(S_{t+1} = s', R_{t+1} = r | S_t = s, A_t = a)$  is a probability that action  $a \in \mathcal{A}$  will lead from state  $s \in \mathcal{S}$  to  $s' \in \mathcal{S}$ , producing a reward  $r \in \mathbb{R}$ ,
- $\gamma \in [0, 1]$  is a discount factor.

# Reinforcement Learning




A Markov decision process is a quadruple  $(\mathcal{S}, \mathcal{A}, P, \gamma)$ , where:

- $\mathcal{S}$  is a set of states,
- $\mathcal{A}$  is a set of actions,
- $P(S_{t+1} = s', R_{t+1} = r | S_t = s, A_t = a)$  is a probability that action  $a \in \mathcal{A}$  will lead from state  $s \in \mathcal{S}$  to  $s' \in \mathcal{S}$ , producing a reward  $r \in \mathbb{R}$ ,
- $\gamma \in [0, 1]$  is a discount factor.

Let a return  $G_t$  be  $G_t \stackrel{\text{def}}{=} \sum_{k=0}^{\infty} \gamma^k R_{t+1+k}$ .

# K-armed Bandits





# K-armed Bandits



Let  $q_*(a)$  be the real *value* of an action  $a$ :

$$q_*(a) = \mathbb{E}[R_{t+1} | A_t = a].$$

# K-armed Bandits



Let  $q_*(a)$  be the real *value* of an action  $a$ :

$$q_*(a) = \mathbb{E}[R_{t+1} | A_t = a].$$

Denoting  $Q_t(a)$  our estimated value of action  $a$  at time  $t$ , we would like  $Q_t(a)$  to converge to  $q_*(a)$ .

# K-armed Bandits



Let  $q_*(a)$  be the real *value* of an action  $a$ :

$$q_*(a) = \mathbb{E}[R_{t+1} | A_t = a].$$

Denoting  $Q_t(a)$  our estimated value of action  $a$  at time  $t$ , we would like  $Q_t(a)$  to converge to  $q_*(a)$ .

A natural way to estimate  $Q_t(a)$  is

$$Q_t(a) \stackrel{\text{def}}{=} \frac{\text{sum of rewards when action } a \text{ is taken}}{\text{number of times action } a \text{ was taken}}.$$

# K-armed Bandits



Following the definition of  $Q_t(a)$ , we could choose a *greedy action*  $A_t$  as

$$A_t(a) \stackrel{\text{def}}{=} \arg \max_a Q_t(a).$$

# K-armed Bandits



Following the definition of  $Q_t(a)$ , we could choose a *greedy action*  $A_t$  as

$$A_t(a) \stackrel{\text{def}}{=} \arg \max_a Q_t(a).$$

## Exploitation versus Exploration

Choosing a greedy action is *exploitation* of current estimates. We however also need to *explore* the space of actions to improve our estimates.

# K-armed Bandits



Following the definition of  $Q_t(a)$ , we could choose a *greedy action*  $A_t$  as


$$A_t(a) \stackrel{\text{def}}{=} \arg \max_a Q_t(a).$$

## Exploitation versus Exploration

Choosing a greedy action is *exploitation* of current estimates. We however also need to *explore* the space of actions to improve our estimates.

An  $\varepsilon$ -*greedy* method follows the greedy action with probability  $1 - \varepsilon$ , and chooses a uniformly random action with probability  $\varepsilon$ .

# K-armed Bandits



# K-armed Bandits

## Incremental Implementation

Let  $Q_n$  be an estimate using  $n$  rewards  $R_1, \dots, R_n$ .

$$\begin{aligned} Q_n &= \frac{1}{n} \sum_{i=1}^n R_i \\ &= \frac{1}{n} \left( R_n + \frac{n-1}{n-1} \sum_{i=1}^{n-1} R_i \right) \\ &= \frac{1}{n} (R_n + (n-1)Q_{n-1}) \\ &= \frac{1}{n} (R_n + nQ_{n-1} - Q_{n-1}) \\ &= Q_{n-1} + \frac{1}{n} (R_n - Q_{n-1}) \end{aligned}$$



# K-armed Bandits



## Non-stationary Problems

Analogously to the solution obtained for a stationary problem, we consider

$$Q_{n+1} = Q_n + \alpha(R_{n+1} - Q_n).$$

# Policies



A *policy*  $\pi$  computes a distribution of actions in a given state, i.e.,  $\pi(a|s)$  corresponds to a probability of performing an action  $a$  in state  $s$ .

# Policies



A *policy*  $\pi$  computes a distribution of actions in a given state, i.e.,  $\pi(a|s)$  corresponds to a probability of performing an action  $a$  in state  $s$ .

## Value Function

To evaluate a quality of policy, we define *value function*  $v_\pi(s)$ , or more explicitly *state-value function*, as

$$v_\pi(s) \stackrel{\text{def}}{=} \mathbb{E}_\pi[G_t | S_t = s].$$

# Policies



A *policy*  $\pi$  computes a distribution of actions in a given state, i.e.,  $\pi(a|s)$  corresponds to a probability of performing an action  $a$  in state  $s$ .

## Value Function

To evaluate a quality of policy, we define *value function*  $v_\pi(s)$ , or more explicitly *state-value function*, as

$$v_\pi(s) \stackrel{\text{def}}{=} \mathbb{E}_\pi[G_t | S_t = s].$$

An *action-value function* for policy  $\pi$  is defined analogously as

$$q_\pi(s, a) \stackrel{\text{def}}{=} \mathbb{E}_\pi[G_t | S_t = s, A_t = a].$$

# Policies



A *policy*  $\pi$  computes a distribution of actions in a given state, i.e.,  $\pi(a|s)$  corresponds to a probability of performing an action  $a$  in state  $s$ .

## Value Function

To evaluate a quality of policy, we define *value function*  $v_\pi(s)$ , or more explicitly *state-value function*, as

$$v_\pi(s) \stackrel{\text{def}}{=} \mathbb{E}_\pi[G_t | S_t = s].$$

An *action-value function* for policy  $\pi$  is defined analogously as

$$q_\pi(s, a) \stackrel{\text{def}}{=} \mathbb{E}_\pi[G_t | S_t = s, A_t = a].$$

It follows that

$$q_\pi(s, a) = \mathbb{E}_\pi[R_{t+1} + \gamma v_\pi(S_{t+1}) | S_t = s, A_t = a].$$

# Optimal Policy



As value functions define a partial ordering of policies ( $\pi' \geq \pi$  if and only if for all states  $s$ ,  $v_{\pi'}(s) \geq v_{\pi}(s)$ ), it can be proven that there always exists an *optimal policy*  $\pi_*$ , which is better or equal to all other policies.

Intuitively,  $\pi_*(s) = \arg \max_a q_*(s, a)$ .

# Optimal Policy



As value functions define a partial ordering of policies ( $\pi' \geq \pi$  if and only if for all states  $s$ ,  $v_{\pi'}(s) \geq v_{\pi}(s)$ ), it can be proven that there always exists an *optimal policy*  $\pi_*$ , which is better or equal to all other policies.

Intuitively,  $\pi_*(s) = \arg \max_a q_*(s, a)$ .

## Policy Improvement Theorem

Let  $\pi$  and  $\pi'$  be any pair of deterministic policies, such that  $q_{\pi}(s, \pi'(s)) \geq v_{\pi}(s)$ . Then  $\pi' \geq \pi$ , i.e., for all states  $s$ ,  $v_{\pi'}(s) \geq v_{\pi}(s)$ .

# Monte Carlo Control



On-policy first-visit MC control (for  $\varepsilon$ -soft policies), estimates  $\pi \approx \pi_*$

Algorithm parameter: small  $\varepsilon > 0$

Initialize:

$\pi \leftarrow$  an arbitrary  $\varepsilon$ -soft policy

$Q(s, a) \in \mathbb{R}$  (arbitrarily), for all  $s \in \mathcal{S}$ ,  $a \in \mathcal{A}(s)$

$Returns(s, a) \leftarrow$  empty list, for all  $s \in \mathcal{S}$ ,  $a \in \mathcal{A}(s)$

Repeat forever (for each episode):

Generate an episode following  $\pi$ :  $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

Loop for each step of episode,  $t = T-1, T-2, \dots, 0$ :

$G \leftarrow G + R_{t+1}$

Unless the pair  $S_t, A_t$  appears in  $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$ :

Append  $G$  to  $Returns(S_t, A_t)$

$Q(S_t, A_t) \leftarrow \text{average}(Returns(S_t, A_t))$

$A^* \leftarrow \arg \max_a Q(S_t, a)$  (with ties broken arbitrarily)

For all  $a \in \mathcal{A}(S_t)$ :

$$\pi(a|S_t) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \varepsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$$



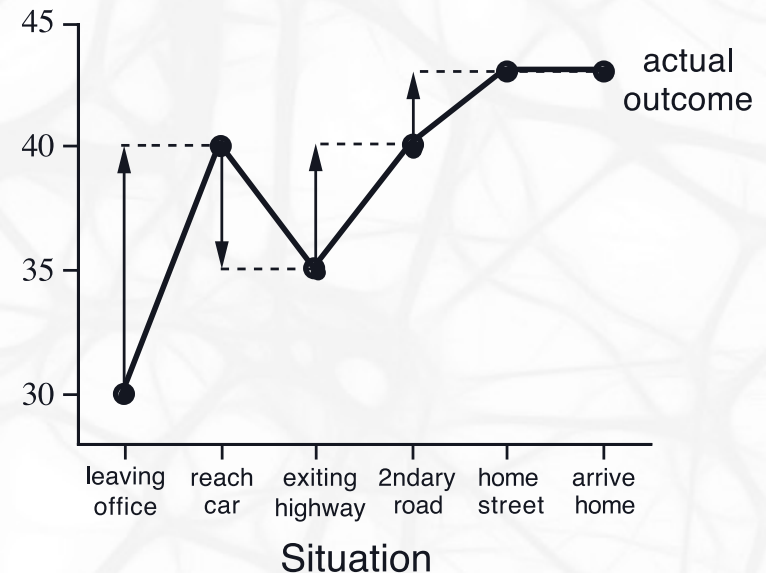

# Temporal Difference Methods



<i>State</i>	<i>Elapsed Time (minutes)</i>	<i>Predicted Time to Go</i>	<i>Predicted Total Time</i>
leaving office, friday at 6	0	30	30
reach car, raining	5	35	40
exiting highway	20	15	35
2ndary road, behind truck	30	10	40
entering home street	40	3	43
arrive home	43	0	43

# Temporal Difference Methods

<i>State</i>	<i>Elapsed Time (minutes)</i>	<i>Predicted Time to Go</i>	<i>Predicted Total Time</i>
leaving office, friday at 6	0	30	30
reach car, raining	5	35	40
exiting highway	20	15	35
2ndary road, behind truck	30	10	40
entering home street	40	3	43
arrive home	43	0	43



# Temporal Difference Methods



A straightforward modification of Monte Carlo algorithm with constant-step update and temporal difference is given by

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha[R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]$$

and is called *Sarsa* ( $S_t, A_t, R_{t+1}, S_{t+1}, A_{t+1}$ ).

# Temporal Difference Methods



A straightforward modification of Monte Carlo algorithm with constant-step update and temporal difference is given by

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha[R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]$$

and is called *Sarsa* ( $S_t, A_t, R_{t+1}, S_{t+1}, A_{t+1}$ ).

**Sarsa (on-policy TD control) for estimating  $Q \approx q_*$**

Algorithm parameters: step size  $\alpha \in (0, 1]$ , small  $\varepsilon > 0$

Initialize  $Q(s, a)$ , for all  $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$ , arbitrarily except that  $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

    Initialize  $S$

    Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

    Loop for each step of episode:

        Take action  $A$ , observe  $R, S'$

        Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

    until  $S$  is terminal

# Q-learning



*Q-learning* is another TD control algorithm by (Watkins, 1989), defined by

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)].$$

# Q-learning

Q-learning is another TD control algorithm by (Watkins, 1989), defined by

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)].$$

**Q-learning (off-policy TD control) for estimating  $\pi \approx \pi_*$**

Algorithm parameters: step size  $\alpha \in (0, 1]$ , small  $\varepsilon > 0$

Initialize  $Q(s, a)$ , for all  $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$ , arbitrarily except that  $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

    Initialize  $S$

    Loop for each step of episode:

        Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)



        Take action  $A$ , observe  $R, S'$

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$


$S \leftarrow S'$

    until  $S$  is terminal

# Sarsa vs Q-learning



# Double Q-learning





# Double Q-learning

## Double Q-learning, for estimating $Q_1 \approx Q_2 \approx q_*$

Algorithm parameters: step size  $\alpha \in (0, 1]$ , small  $\varepsilon > 0$

Initialize  $Q_1(s, a)$  and  $Q_2(s, a)$ , for all  $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$ , such that  $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

Initialize  $S$

Loop for each step of episode:

Choose  $A$  from  $S$  using the policy  $\varepsilon$ -greedy in  $Q_1 + Q_2$

Take action  $A$ , observe  $R, S'$

With 0.5 probability:

$$Q_1(S, A) \leftarrow Q_1(S, A) + \alpha \left( R + \gamma Q_2(S', \arg\max_a Q_1(S', a)) - Q_1(S, A) \right)$$

else:

$$Q_2(S, A) \leftarrow Q_2(S, A) + \alpha \left( R + \gamma Q_1(S', \arg\max_a Q_2(S', a)) - Q_2(S, A) \right)$$

$S \leftarrow S'$

until  $S$  is terminal